

1. THE BRAND RULES, AND HOW I DERIVED THEM

Four layers, all data in `config/`, so a rule change is a row and not a deploy.

Voice. From the site and the packaging.

The positioning is not "clean skincare", it is that the concentration is printed on the front of the pack.

So: state the active and its exact strength, let the number be the proof.

Compliance. Six rule families, severity-tagged, each with a written rationale.

- *ASCI code.* Requires substantiation on demand for any objectively ascertainable claim. Kills "clinically proven product", leaves "clinically studied ingredient" standing.
- *Drugs and Cosmetics Rules 1945.* A cosmetic claiming to treat a disease becomes a drug.

The rationale matters more than the patterns: it lets a reviewer argue with the rule, not the score.

Creative grammar. The layer I would defend hardest, because it changed the output and came from iteration rather than authorship.

My first visual rules were written from the brand's stated aesthetic. Every generation came back a tasteful bottle on warm sand: every constraint satisfied, nothing like the brand.

I did not fix that by rewriting instructions in the abstract. I put real generations beside three banners Minimalist is currently running and rewrote the rules against the gap:

- Flat white, not warm sand
- The product on white geometric pedestals
- Hairline leader lines to small-caps labels

Closer to a spec sheet than a photograph. That became four named archetypes and six devices, and the brief model must *name* one rather than describe a layout, so the image prompt expands its construction rules deterministically.

Platform. A property of every placement, not a preset.

A 40-character headline that fits Meta is cut by Google at 30, and mixing them writes copy to the tighter limit.

These are my reading of the brand, not its guidelines.

2. WHAT I CUT, AND WHY

Performance prediction. No outcome data exists, so a predicted CTR is a confident number with nothing behind it: the failure this brief warns about, dressed as a feature.

The reviewer workflow (comments, sign-off, Slack). The stated pain and the hardest cut. A workflow around a bad standard routes bad decisions faster.

Firecrawl. Started with to dodge bot protection, then removed entirely. Removing it fixed a bug: enrichment was Firecrawl-only and returned "no page, no warning" on failure, so a brief ran on one line of description and nothing said so. A thin brief and a model ignoring its brief look identical from outside. Shopify serves the ingredients and FAQ in the HTML anyway.

Video, multi-variant testing, international compliance. Scope.

Deliberation on three things I took rather than decided, because calling them evaluations would be invented rigour.

CHOICE	WHY
Nano Banana Pro	Picked on reputation. Quality was the constraint, cost was not, it worked, I stopped looking.
Railway	Collapses hosting, Postgres and secrets into one place.
One password	Covers everything, because every action behind it spends a live key.

One note on the brief. It calls the delay reviewer disagreement, then asks for a scorer. Those do not meet. Scoring makes one reviewer faster; disagreement is a question of who decides. So the scorer is loud on compliance, advisory on taste.

3. THE DECISION I WAS LEAST SURE ABOUT

Whether the image model should render the pack at all.

It is given the real photograph as a reference, so the packaging is reproduced rather than invented. Reproduced is not composited. The model redraws the bottle, and on this brand the printed concentration is the claim.

A regenerated label reading 12% on a 10% pack is a false regulated claim, and it passes every gate I built, because the claim gate reads the copy layer and not the pixels of the bottle.

The safe answer is compositing the real photograph into an HTML or SVG layout. I did not, because this brand's real ads are constructed rather than photographed, and a composite pipeline needs a layout engine per archetype per aspect ratio, which is where the whole build would have gone.

So I moved the risk rather than removing it: the claim gate runs on text, before any image is paid for, so the words are verified free. The pixels are not.

The fix is OCR on the render against the product data. **I did not build it, and I would not put spend behind this until it exists.**